

## Retrieval Augmented Generation Overview

RAG combines a retriever with a generator model.

The retriever finds relevant chunks from a vector store.

The generator conditions its answer on those chunks.

## Chunking and Embeddings

Documents are split into overlapping chunks before embedding.

Each chunk becomes a dense vector stored in ChromaDB.

Overlap prevents context loss at chunk boundaries.